



Certified Kubernetes Administrator Prep

Volumes and Their Access Modes

Volumes and Their Access Modes

- PersistentVolume -- API for users that abstracts implementation details of storage
- PersistentVolumeClaim -- Method for users to claim durable storage regardless of implementation details



PersistentVolume (PV)

- Provisioned storage in the cluster
- Cluster resource
- Volume plugins (as discussed in the previous lesson) have independent lifecycle from pods
- Volumes share the lifecycle of the pod; PersistentVolumes do **not**
- API object (YAML) details the implementation



PersistentVolumeClaim (PVC)

- Request for storage
- Pods consume node resources; PVCs consume PV resources.
- Pods can request specific CPU and memory; PVCs can request specific size and access modes.



PVs and PVCs

- Users and applications do not share identical requirements
- Administrators should offer a variety of PVs without users worrying about the implementation details
- PVs are cluster resources, PVCs are requests for the cluster resource
- PVCs also act as a “claim check” on a resource
- PVs and PVCs have a set lifecycle
 - Provision
 - Bind
 - Reclaim



Provisioning

- **Static**
 - Creates PVs
 - In the K8S API and available for consumption
- **Dynamic**
 - Used when none of the static PVs match the PVC
 - Based on StorageClasses
 - PVC must request a created and configured storage class
 - Claims requesting nameless class disable dynamic provisioning

To enable dynamic storage provisioning, `DefaultStorageClass` admission controller on the API server must be enabled.



Binding

- User creates PVC
- Master watches for new PVCs and matches them to PVs
- Master binds PVC to PV
- Volume may be more than the request
- Binds are exclusive
- PVC -> PV mapping is always 1:1
- Claims not matched will remain unbound indefinitely



Pod Uses Volume

- Pods treat PVCs as volumes
- Cluster checks claim, mounts appropriate volume to pod



Persistent Volume Claim Protection

- Alpha feature as of K8s 1.9
- Ensures PVCs actively in use do not get removed from the system
- PVC considered active when:
 - The pod status is Pending and the pod is assigned to a node
 - The pod status is Running
- If a user deletes a PVC in use, removal is postponed until PVC is not in use by any pod



Persistent Volume Claim Protection

```
$> kubectl describe pvc hostpath
```

```
Name:          hostpath
```

```
Namespace:     default
```

```
StorageClass:  example-hostpath
```

```
Status:        Terminating
```

```
Volume:
```

```
Labels:        <none>
```

```
Annotations:   volume.beta.kubernetes.io/storage-class=example-hostpath
```

```
volume.beta.kubernetes.io/storage-provisioner=example.com/hostpath
```

```
Finalizers:    [kubernetes.io/pvc-protection]
```



Reclaiming

- User can delete PVC objects
- Reclaim policy for a PV tells the cluster what to do with the volume
- Policies:
 - Retain
 - Recycle
 - Delete
- Policies allow for manual reclamation of a resource.
 - PVC deleted
 - PV still exists; volume is “released”
 - Not yet available because the data is still present
 - Admin can manually reclaim the volume



Reclaiming

- Administrator can configure custom recycler pod template
- Must contain a volumes specification
- Volume plugs that support Delete reclaim policy
 - Removes the PersistentVolume object from K8s
 - Associated storage assets in the infrastructure
 - Dynamically provisioned volumes inherit reclaim policy of their StorageClass
- Administrator should configure StorageClass according to expectations



Capacity

- PVs have a specific storage capacity
- Set using “capacity” attribute
- Storage size is the only resource that can be set or requested
- Future attribute plans:
 - IOPS
 - Throughput
 - ???
- Prior to K8s v1.9, default behavior for all volume plugins was to create a filesystem
- As of 1.9, however, user can specify volumeMode
 - Raw Block Devices -- “Block”
 - File systems -- “Filesystem” (default)



Access Modes

- Must be supported by storage resource provider
- ReadWriteOnce -- Can be mounted as read/write by one node only (RWO)
- ReadOnlyMany -- Can be mounted read-only by many nodes (ROX)
- ReadWriteMany -- Can be mounted read/write by many nodes (RWX)

A volume can only be mounted using one access mode at a time, regardless of the modes that are supported.



Conclusion

- PVC vs PV
- Storage object lifecycle
- Access modes

