



Linux Academy
Live! Lab

Auto-Scaling
Based Off
SQS Message
Queue Size

Contents

Introduction.....	1
Creating the SQS Queue.....	1
Sending a Message.....	1
Scaling with CloudWatch.....	2
Configuring EC2 Launch Configurations.....	3
Testing the Scale.....	4

Lab Connection Information

- Labs may take up to five minutes to build
- Access to an AWS Console is provided on the Live! Lab page, along with your login credentials
- Ensure you are using the N. Virginia region
- Labs will automatically end once the allotted amount of time finishes

Related Courses

[*AWS CSA - Pro Course*](#)

Related Videos

[*SQS Message Priority*](#)

[*SQS Batch Processing Job Observer*](#)

Need Help?

[*Linux Academy Community*](#)

*... and you can
always send in a
support ticket on
our website to talk
to an instructor!*

Introduction

In this lab, we learn how to auto-scale your infrastructure through a combination of *CloudWatch*, *SQS* messages, and AWS's *Auto Scaling* feature. First, we create an SQS queue, and then, through CloudWatch, set up an alarm to trigger every time we hit 40 messages in our queue. When this alarm triggers, another instance is added.

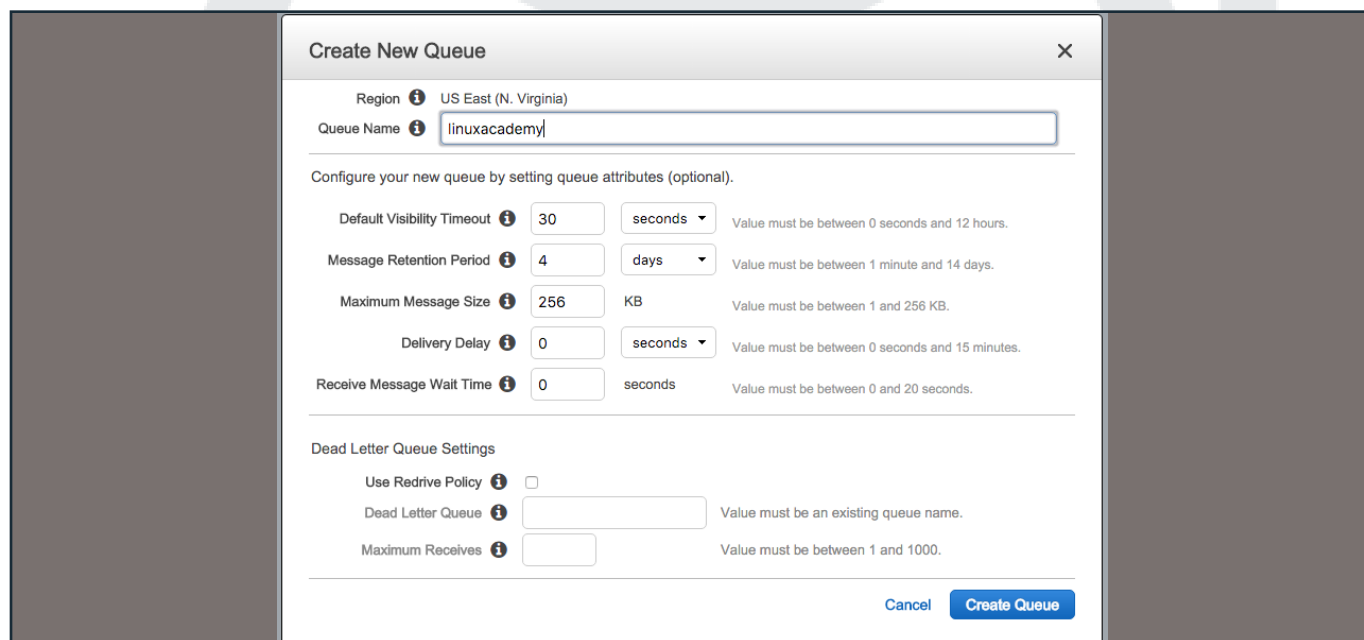
Log in to your **AWS Console** with the credentials provided on the **Live! Lab** page.

Within your **EC2 Dashboard**, you can see an instance called *Ops* — this is the instance we are using to send messages to our SQS queue.

However, we first need to create the queue.

Creating the SQS Queue

Navigate to the **SQS Dashboard** and press **Get Started Now**.



Enter a **Queue Name**. We choose *linuxacademy* for our queue. We left the *default options* for the rest of the queue settings.

Press **Create Queue**.

Sending a Message

To send an initial message to our new SQS queue, we need to log in to

For information regarding these attributes, review the SQS videos for the [AWS Certified Developer - Associate](#) course.

our *Ops* instance. Open your terminal on your workstation. You can retrieve the *IP address* of your instance from the **EC2 Dashboard**. The username you are logging in with is *linuxacademy* with a password of *123456*.

su into the root user. The password is also *123456*.

Now configure the AWS Cli:

```
aws configure
```

You are asked for an **AWS Access Key ID**, press *enter* to select the default, *none*. Do the same for the **AWS Secret Access Key**. The **Default region name** should be set to *us-east-1*. The **Default output format** should also be left to *none*.

To send a message to our SQS queue, we need the URL to that queue. From the **SQS Dashboard**, select your *linuxacademy* queue, and *copy the URL provided*. The URL is located in the **Details** section.

Send a message to the queue from the command line, replacing *https://sqs.us-east-1.amazonaws.com/642132563443/linuxacademy* with *your queue's URL*:

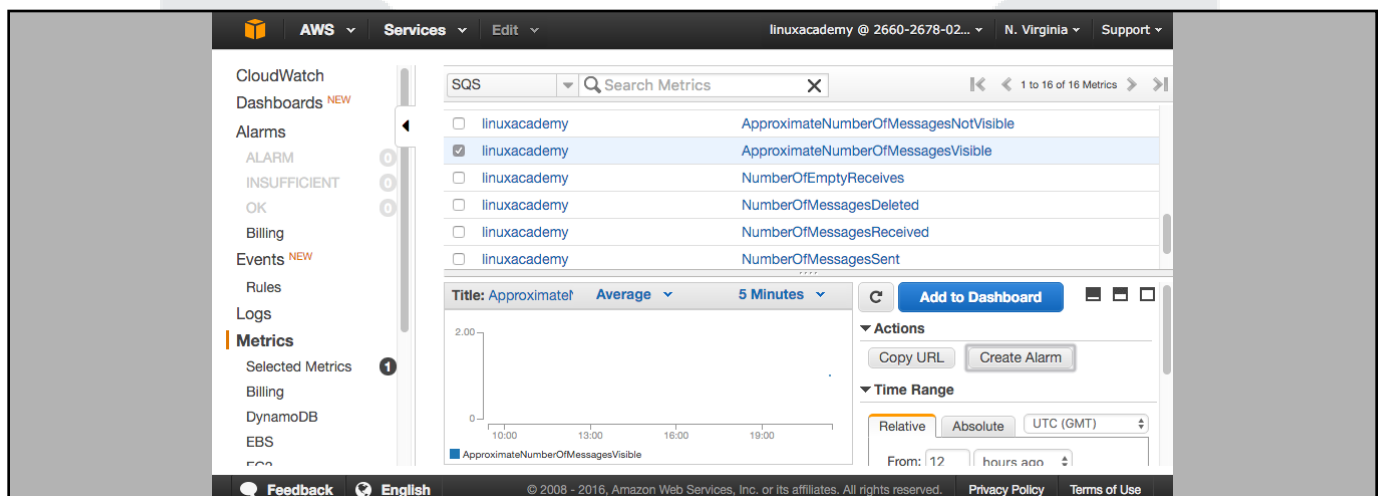
```
aws sqs send-message --queue-url https://sqs.us-east-1.amazonaws.com/642132563443/linuxacademy --message-body linuxacademytest
```

The *message-body* can be any form of example text.

If you return to/refresh your queue, you should now see *1 Message[s] Available*.

Scaling with CloudWatch

Through the use of CloudWatch, we can configure our infrastructure to scale based on how many messages are in our queue. From the **CloudWatch Dashboard**, select **SQS** from the left menu.



We are working with the *linuxacademy* queue (under **QueueName**) — if it is not yet available, refresh the page until it is.

We want to create an alarm based on the **ApproximateNumberOfMessagesVisible** metric. Select this, and press **Create Alarm**.

For **Name** and **Description** we put *sqs-messages-in-queue*.

We want CloudWatch to set an alarm when we hit a certain number of messages. This alarm deploys out new instances based on the auto-scaling policy we use.

Set the **Whenever: ApproximateNumberOfMessagesVisible** to ≥ 40 , for a period of *1 Minute*. Delete the notification option. Select **Create Alarm**.

Configuring EC2 Launch Configurations

From the **EC2 Dashboard**, select **Launch Configurations**, **Create Auto Scaling group**. Press **Create Launch Configuration**.

Select the **Amazon Linux AMI** and the *General purpose t2.micro*. Select **Next: Configure details**.

We gave our configuration a **Name** of *sqs-pool-workers*, and selected **Next: Add Storage**. Leave all storage settings as-is, and press **Next: Configure security group**. This can also be left as-is. Press **Review, Create**

launch configuration. A pop-up asks you to select a key pair; from the initial drop-down, select *proceed without a key pair*. **Create launch configuration.**

The screenshot shows the AWS Management Console interface for creating an Auto Scaling Group. The top navigation bar includes the AWS logo, 'Services', 'Edit', and user information. The main content area has a progress bar with five steps: 1. Configure Auto Scaling group details (active), 2. Configure scaling policies, 3. Configure Notifications, 4. Configure Tags, and 5. Review. The 'Create Auto Scaling Group' title is at the top of the main section, with a 'Cancel and Exit' link. Below the title are two error messages in red boxes: 'No default VPC found' and 'Unable to retrieve load balancers: User: arn:aws:iam::266026780299:user/linuxacademy is not authorized to perform: elasticloadbalancing:DescribeLoadBalancers'. The 'Launch Configuration' section includes a dropdown for 'Launch Configuration' (sqs-pool-workers), a text field for 'Group name' (sqs-pool-as-group), a dropdown for 'Group size' (Start with 1 instances), a dropdown for 'Network' (vpc-0faca76b (10.0.0.0/16)) with a 'Create new VPC' button, and a dropdown for 'Subnet' (subnet-188b0c32 (10.0.2.0/24) | us-east-1a and subnet-e51b6d93 (10.0.0.0/24) | us-east-1b) with a 'Create new subnet' button. A yellow warning box states 'No public IP addresses will be assigned'. At the bottom right, there are 'Cancel' and 'Next: Configure scaling policies' buttons.

We now need to create an Auto Scaling group. Our **Group** name is set to *sqs-pool-as-group*. Under **Network**, select the single available VPC, and then choose both subnets for the **Subnet** values. Press **Next: Configuring scaling policies**.

We want to *Use scaling policies to adjust the capacity of this group* to scale between *1* and *4 instances*.

For **Execute policy when**, select your *sqs-messages-in-queue* alarm, and set **Take the action** for *1 instance*. Our **Instances need** is set to *10 seconds*.

Because this is a lab environment, we are working only with scaling up. Remove the decrease policy below, or use your new skills to try to set it up yourself.

Press **Review, Create Auto Scaling group**.

Testing the Scale

To test our new policies, we need to quickly generate forty or more messages for our queue. Return to your terminal, where you should still be logged into the *Ops* instance as *root*.

Using your text editor of choice, create a script called *sqs.sh*, replacing the *queue-url* appropriately:

```
#!/bin/bash
for i in $(seq 1 60)
do
    aws sqs send-message --queue-url https://sqs.us-east-1.amazonaws.com/642132563443/
    linuxacademy --message-body linuxacademytest
done
```

Make the script executable:

```
chmod +x sqs.sh
```

Finally, run the script:

```
./sqs.sh
```

If you return to your **SQS queue**, you can refresh to see your number of **Messages Available**. When the number reaches **61**, it is finished.

Navigate to your **EC2 Dashboard** to confirm more instances are spinning up.